

WEB INTELLIGENCE

Lecture 6: Language Modeling II

Russa Biswas

October 21, 2025



AALBORG UNIVERSITY

(Based on slides of Dan Jurafsky, Andrew Ng)



Language Modeling - Recap

Probability of a sentence = Probability of a sequence of the words in the sentence

$$P(S) = P(w_1 w_2 \dots w_n)$$



Chain Rule

$$P(w_1 w_2 \dots w_n) = \prod_i P(w_i | w_1 w_2 \dots w_{i-1})$$



Markov Assumption

$$P(w_i | w_1 w_2 \dots w_{i-1}) \approx P(w_i | w_{i-1}) \quad \text{N-gram model}$$



$$P(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})} \quad \text{Maximum Likelihood Estimation}$$

How to evaluate a Language Model?

Two ways to evaluate: **Extrinsic Evaluation** and **Intrinsic Evaluation**

- **Extrinsic evaluation** captures how useful the model is in a particular task (also known as **in-vivo** evaluation)
 - To compare models A and B
 - Put each model in a real task (such as Machine Translation, speech recognition, etc.)
 - Run the task, get a score for A and for B (such as How many words translated correctly, How many words transcribed correctly etc.)
 - Compare accuracy for A and B
- **Limitations:**
 - Expensive, time-consuming
 - Doesn't always generalize to other applications

How to evaluate a Language Model?

Intrinsic evaluation captures how well the model captures what it is supposed to capture (also known as in-vitro evaluation): ***Perplexity***

- **Perplexity**
 - Directly measures language model performance at predicting words.
 - Doesn't necessarily correspond with real application performance
 - But gives us a single general metric for language models
 - Useful for large language models (LLMs) as well as n-grams

How to evaluate a Language Model?

Divide the data/corpus into **training** set and **test** set (disjoint sets)

We train parameters of our model on a **training set**.

We test the model's performance on data we haven't seen.

- A **test set** is an unseen dataset; different from training set.
 - Intuition: we want to measure generalization to unseen data
 - An **evaluation metric** (like **perplexity**) tells us how well our model does on the test set.
- If we're building an LM for a specific task
 - The test set should reflect the task language we want to use the model for (domain specific – medical, legal, etc.)
- If we're building a general-purpose model
 - We'll need lots of different kinds of training data (general data covering all domains)
 - We don't want the training set or the test set to be just from one domain or author or language.

How to evaluate a Language Model?

- ***Training on the Test set***: If sentences from the test set is present on the training data
 - Issues:
 - LM will assign that sentence an artificially high probability when we see it in the test set
 - hence assign the whole test set a falsely high probability.
 - the LM would look better than it really is
 - Sentences from test set **should not be** in the training set
 - Dev Set:
 - If we test on the test set many times we might implicitly tune to its characteristics (overfit)
 - Noticing which changes make the model better.
 - Therefore, we need a third dataset: **development test set** or **devset** or **validation set**.
 - We test our LM on the devset until the very end to adjust the LM
 - And then test our LM on the **test set** once

Perplexity

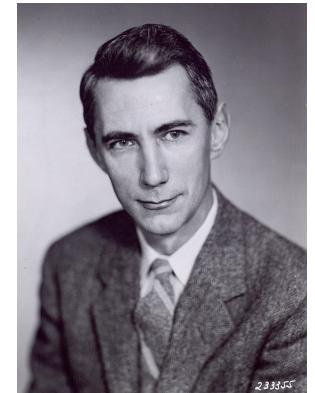
Intuition: A good LM prefers "real" sentences

- Assign higher probability to "real" or "frequently observed" sentences
- Assigns lower probability to "word salad" or "rarely observed" sentences?

The Shannon Game: **How well can we predict the next word?**

- Once upon a _____
- That is a picture of a _____

time	0.9
dream	0.03
midnight	0.02
...	
and	1e-100



Claude Shannon

- Unigrams are terrible at this game (Why?)
- A good LM is one that assigns a higher probability to the next word that actually occurs

Picture credit: Historiska bildsamlingen
<https://creativecommons.org/licenses/by/2.0/>

Perplexity

The best LM assigns high probability to the entire test set

When comparing two LMs, A and B

- We compute $P_A(\text{test set})$ and $P_B(\text{test set})$
- The better LM will give a higher probability to the test set than the other LM.
- Probability depends on size of test set
 - Probability gets smaller the longer the text
 - Better: a metric that is **per-word**, normalized by length

Perplexity is the inverse probability of the test set,

$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

normalized by the number of words

Perplexity

- The inverse comes from the original definition of perplexity from cross-entropy rate in information theory)
- Probability range is $[0,1]$, perplexity range is $[1,\infty]$
- **Minimizing perplexity is the same as maximizing probability**

$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

Chain rule: $PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$

Bigrams: $PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$

Perplexity – Weighted Average Branching Factor

- The branching factor of a language is the number of possible next words that can follow any word
- The Perplexity (PP) is referred to as the **weighted average branching factor** and is the most common evaluation metric for N-gram language models

Example: Deterministic language $L = \{\text{red, blue, green}\}$

Branching factor = 3 (any word can be followed by red, blue, green)

Now assume LM A where each word follows any other word with **equal probability $\frac{1}{3}$**

Given a test set $T = \text{"red red red red blue"}$

$$\text{Perplexity}_A(T) = P_A(\text{red red red red blue})^{-1/5} = ((\frac{1}{3})^5)^{-1/5} = (\frac{1}{3})^{-1} = 3$$

- But now suppose red was very likely in training set, such that for LM B: $P(\text{red}) = 0.8$ $p(\text{green}) = 0.1$ $p(\text{blue}) = 0.1$
- We would expect the probability to be higher, and hence the perplexity to be smaller:

$$\text{Perplexity}_B(T) = P_B(\text{red red red red blue})^{-1/5} = (0.8 * 0.8 * 0.8 * 0.8 * 0.1)^{-1/5} = 0.04096^{-1/5} = 0.527^{-1} = 1.89$$

Perplexity

Wall Street Journal (19,979 word vocabulary)

- Training set: 38 million words
- Test set: 1.5 million words

	Unigram	Bigram	Trigram
Perplexity	962	170	109

Lower perplexity = better language model

Since language model probabilities are very small, multiplying them together often yields to underflow. It is often better to use logarithms instead, so

$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$



$$PP(S) = \exp \left(-\frac{1}{n} \sum_{i=1}^n \log_e P(w_i | w_1, \dots, w_{i-1}) \right)$$

Evaluation: Accuracy

How to objectively measure the quality of a (classification) experiment?

- Compare your achieved results with a ground truth (gold standard).
- How to achieve a ground truth?
 - Crowd sourcing
- How to compare achieved results with a ground truth?

Scenario:

Binary Classification task: Goal is to predict whether an email is **Spam** (Positive) or **Not Spam** (Negative).

Train set: 10,000 emails, test set: 100 emails and the results are as follows:

Evaluation: Accuracy

Scenario:

Binary Classification task: Goal is to predict whether an email is **Spam** (Positive) or **Not Spam** (Negative).

Train set: 10,000 emails, test set: 100 emails and the results are as follows:

The number of spam emails correctly predicted as spam - 40

The number of not spam emails correctly predicted as not spam - 45

The number of not spam emails incorrectly predicted as spam - 5 Type I error

The number of spam emails incorrectly predicted as not spam - 10 Type II error

	Predicted: Spam	Predicted: Not Spam
Actual/Ground Truth: Spam	40 True Positives (TP)	10 False Negatives (FN)
Actual/Ground Truth: Not Spam	5 False Positives (FP)	45 True Negatives (TN)

Evaluation: Accuracy

	Predicted: Spam	Predicted: Not Spam
Actual/Ground Truth: Spam	40 True Positives (TP)	10 False Negatives (FN)
Actual/Ground Truth: Not Spam	5 False Positives (FP)	45 True Negatives (TN)

Accuracy: Proportion of correctly classified instances out of all instances.

$$\begin{aligned} \text{Accuracy} &= \frac{TP + TN}{TP + TN + FP + FN} \\ &= \frac{40 + 45}{40 + 10 + 5 + 45} = \frac{85}{100} = 0.85 \end{aligned}$$

Evaluation: Precision and Recall

	Predicted: Spam	Predicted: Not Spam
Actual/Ground Truth: Spam	40 True Positives (TP)	10 False Negatives (FN)
Actual/Ground Truth: Not Spam	5 False Positives (FP)	45 True Negatives (TN)

Precision is the **fraction of retrieved instances that are relevant**

$$precision = \frac{true\ positive}{true\ positive + false\ positive}$$

$$Precision = \frac{40}{40+45} = \frac{40}{85} = 0.89$$

Recall is the **fraction of relevant instances that are retrieved/predicted**

$$recall = \frac{true\ positive}{true\ positive + false\ negative}$$

$$Recall = \frac{40}{40 + 10} = \frac{40}{50} = 0.80$$

Evaluation: F1 score

	Predicted: Spam	Predicted: Not Spam
Actual/Ground Truth: Spam	40 True Positives (TP)	10 False Negatives (FN)
Actual/Ground Truth: Not Spam	5 False Positives (FP)	45 True Negatives (TN)

F1-score is the harmonic mean of precision and recall, and provides a balance between the two.

$$F_1 = 2 \cdot \frac{precision \cdot recall}{precision + recall}$$

$$= 2 * \frac{0.89 * 0.80}{0.89 + 0.80} = 0.84$$

Evaluation

	Predicted:Positive	Predicted: Negative
Actual/Ground Truth: Positive	30 True Positives (TP)	20 False Negatives (FN)
Actual/Ground Truth: Negative	10 False Positives (FP)	40 True Negatives (TN)

Compute Accuracy, Precision, Recall and F1 score

One-hot encoding

One-hot encoding is a method of representing data as binary vectors

Vocabulary: 10,000 unique words

Size of each word vector = 10,000

$$\mathbf{o}_{man} = [00000 \dots 000 \dots 1 \dots 000 \dots 00000000] = \mathbf{o}_{2345}$$

$$\mathbf{o}_{woman} = [00000 \dots 000 \dots 001 \dots 00000000] = \mathbf{o}_{3679}$$

$$\mathbf{o}_{king} = [00000 \dots 001 \dots 000 \dots 000 \dots 00000000] = \mathbf{o}_{356}$$

$$\mathbf{o}_{queen} = [00000 \dots 000 \dots 000 \dots 001 \dots 00000000] = \mathbf{o}_{1892}$$

$$\mathbf{o}_{orange} = [00000 \dots 000 \dots 1 \dots 000 \dots 00000000] = \mathbf{o}_{489}$$

$$\mathbf{o}_{apple} = [00000 \dots 000 \dots 1 \dots 000 \dots 00000000] = \mathbf{o}_{4432}$$

I want a glass of apple juice

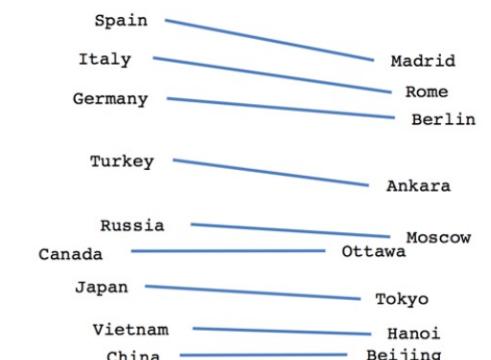
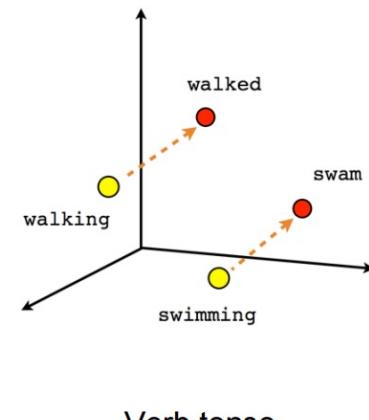
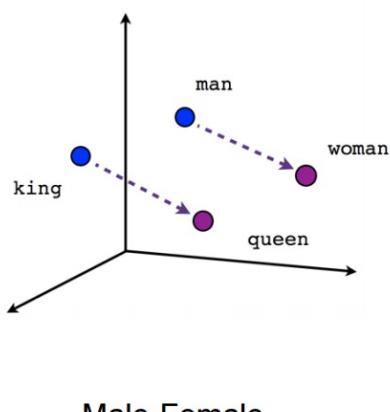
I want a glass of orange _____

One hot encoding

	Man	Woman	King	Queen	Apple	Orange
Gender	-1	1	-0.95	0.95	0.00	0.01
Royal	0.01	0.02	0.93	0.96	-0.01	0.00
Food	-0.001	0.001	-0.001	0.001	0.95	0.98
---size, color, alive, age, cost, etc.						

Features: 300

I want a glass of apple **juice**
 I want a glass of orange **juice**



Embedding Matrix

w1	w2	wn
0.1	-0.4	0.8
-0.5	0.5	-0.1
0.2	-0.9	0.02
...
0.1	0.2	-0.9

$$\mathbf{o}_{orange} = [00000 \dots 000 \dots 1 \dots 000 \dots 00000000] = \mathbf{o}_{489}$$

Size of Matrix E is 300 x 10000

Matrix multiplication of E and O = embedding vector

Size of \mathbf{o}_{orange} is 10000 x 1

Word Embeddings

One way to define "usage": words are defined by their environments (the words around them)

Zellig Harris (1954):

- **If A and B have almost identical environments we say that they are synonyms.**
- We define a word as a vector
- Called an "embedding" because it's embedded into a space: standard way to represent meaning in NLP
- **Every modern NLP algorithm uses embeddings as the representation of word meaning**
- Fine-grained model of meaning for similarity

Word Embeddings – Why vectors?

- Consider sentiment analysis:
 - With **words**, a feature is a word identity
 - Feature 5: 'The previous word was "terrible"
 - requires **exact same word** to be in training and test
 - With **embeddings**:
 - Feature is a word vector
 - 'The previous word was vector [35,22,17...]
 - Now in the test set we might see a similar vector [34,21,14]
 - We can generalize to **similar but unseen words!!!**
- **tf-idf**
 - Information Retrieval!
 - A common baseline model
 - **Sparse** vectors
 - Words are represented by (a simple function of) the **counts** of nearby words
- **Word2vec**
 - **Dense** vectors
 - Representation is created by training a classifier to **predict** whether a word is likely to appear nearby
 - Later we'll discuss extensions called **contextual embeddings**

Word Embeddings

- Sparse versus dense vectors
- tf-idf vectors are
 - **long** (length $|V| = 20,000$ to $50,000$)
 - **sparse** (most elements are zero)
- Alternative: learn vectors which are
 - **short** (length 50-1000)
 - **dense** (most elements are non-zero)
- Why dense vectors?
 - Short vectors may be easier to use as **features** in machine learning (fewer weights to tune)
 - Dense vectors may **generalize** better than explicit counts
- They may do better at capturing synonymy:
 - *car* and *automobile* are synonyms; but are distinct dimensions
 - a word with *car* as a neighbor and a word with *automobile* as a neighbor should be similar, but aren't
- **In practice, they work better**

Word Embeddings

- Neural Language Model"-inspired models
 - Word2vec (skipgram, CBOW), Glove
- Singular Value Decomposition (SVD)
 - A special case of this is called LSA – Latent Semantic Analysis
- Alternative to these "static embeddings":
 - Contextual Embeddings (ELMo, BERT)
 - Compute distinct embeddings for a word in its context
 - Separate embeddings for each token of a word
 - We'll return to this in a later chapter
 - All the new age LLMs

Word2vec

- Instead of **counting** how often each word w occurs near "apricot"
 - Train a classifier on a binary **prediction** task:
 - Is w likely to show up near "apricot"?
- We don't actually care about this task
 - But we'll take the learned classifier weights as the word embeddings
- Big idea: **self-supervision**:
 - A word c that occurs near apricot in the corpus asks as the gold "correct answer" for supervised learning
 - No need for human labels
 - Bengio et al. (2003); Collobert et al. (2011)
- Popular embedding method
- Very fast to train
- Code available on the web
- Idea: **predict** rather than **count**

Word2vec

Word2vec provides a variety of options.

skip-gram with negative sampling (SGNS)

Continuous bag of words approach (CBOW)

Approach: predict if candidate word c is a "neighbor" – Skip-gram

1. Treat the target word t and a neighboring context word c as **positive examples**.
2. Randomly sample other words in the lexicon to get negative examples
3. Use logistic regression to train a classifier to distinguish those two cases
4. Use the learned weights as the embeddings

Word2vec

- Assume a +/- 2 word window, given training sentence:

...lemon, a [tablespoon of apricot jam, a] pinch...

c1 c2 **[target]** c3 c4

Goal: train a classifier that is given a candidate (**word, context**) pair
(apricot, tablespoon)
(apricot, land)

...

And assigns each pair a probability:

$$P(+ | w, c)$$

Word2vec

Remember: two vectors are similar if they have a high dot product

- Cosine is just a normalized dot product

So: $\text{Similarity}(w, c) \propto w \cdot c$

We'll need to normalize to get a probability

- (cosine isn't a probability either)
- **Turning dot products into probabilities**
- $\text{Sim}(w, c) \approx w \cdot c$
- To turn this into a probability
- We'll use the sigmoid from logistic regression:

$$P(+|w, c) = \sigma(c \cdot w) = \frac{1}{1 + \exp(-c \cdot w)}$$

$$\begin{aligned} P(-|w, c) &= 1 - P(+|w, c) \\ &= \sigma(-c \cdot w) = \frac{1}{1 + \exp(c \cdot w)} \end{aligned}$$

Word2vec

- How Skip-Gram Classifier computes $P(+|w, c)$

$$P(+|w, c) = \sigma(c \cdot w) = \frac{1}{1 + \exp(-c \cdot w)}$$

This is for one context word, but we have lots of context words.

We'll assume independence and just multiply them:

$$P(+|w, c_{1:L}) = \prod_{i=1}^L \sigma(c_i \cdot w)$$

$$\log P(+|w, c_{1:L}) = \sum_{i=1}^L \log \sigma(c_i \cdot w)$$

Word2vec

...lemon, a [tablespoon of apricot jam, a] pinch...

- c1 c2 [target] c3 c4

positive examples +

t c

apricot tablespoon

apricot of

apricot jam

apricot a

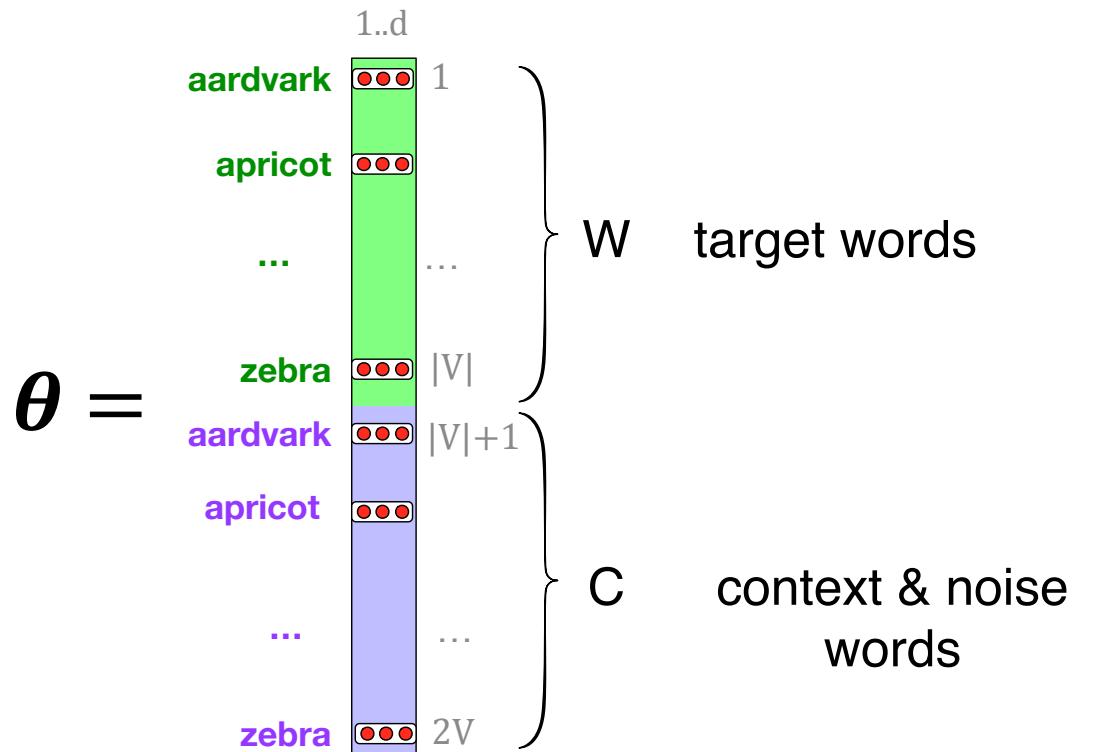
For each positive example we'll grab k negative examples, sampling by frequency

negative examples -

t	c	t	c
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if

Word2vec

- A probabilistic classifier that,
 - given a test target word w
 - its context window of L words $c_{1:L}$,
- assigns a probability that w occurs in this window.
- To compute this, we just need embeddings for all the words.



Word2vec

- Given the set of positive and negative training instances, and an initial set of embedding vectors
- The goal of learning is to adjust those word vectors such that we:
 - **Maximize** the similarity of the **target word, context word** pairs (w, c_{pos}) drawn from the positive data
 - **Minimize** the similarity of the (w, c_{neg}) pairs drawn from the negative data.

Word2vec

Loss function for one w with $c_{pos}, c_{neg1} \dots c_{negk}$

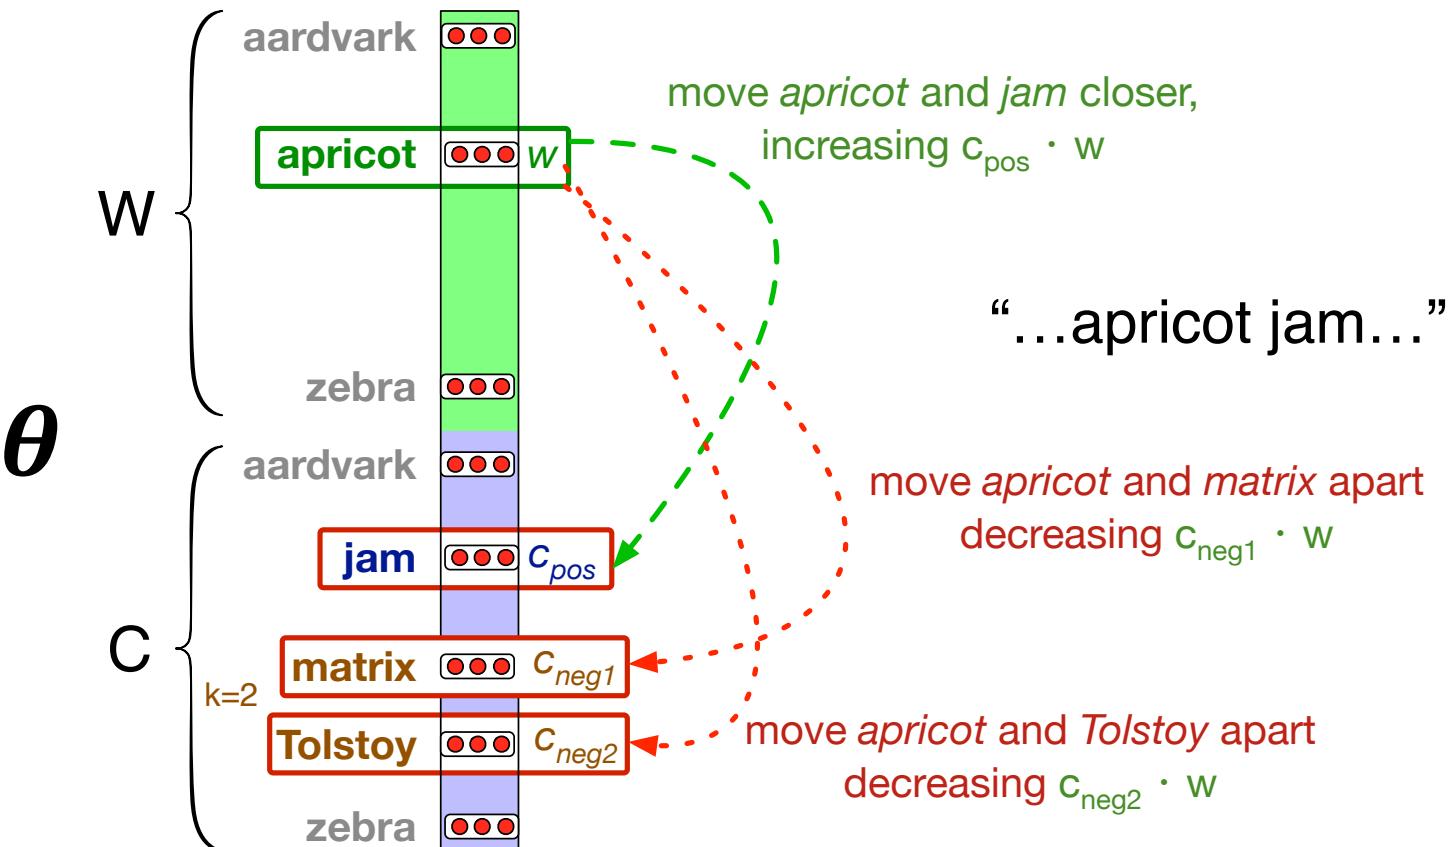
Maximize the dot product of the word with the actual context words, and

minimize the dot products of the word with the k negative sampled non-neighbor words.

$$\begin{aligned} L_{CE} &= -\log \left[P(+|w, c_{pos}) \prod_{i=1}^k P(-|w, c_{neg_i}) \right] \\ &= - \left[\log P(+|w, c_{pos}) + \sum_{i=1}^k \log P(-|w, c_{neg_i}) \right] \\ &= - \left[\log P(+|w, c_{pos}) + \sum_{i=1}^k \log (1 - P(+|w, c_{neg_i})) \right] \\ &= - \left[\log \sigma(c_{pos} \cdot w) + \sum_{i=1}^k \log \sigma(-c_{neg_i} \cdot w) \right] \end{aligned}$$

Word2vec

- How to learn?
 - Stochastic gradient descent!
- We'll adjust the word weights to
 - make the positive pairs more likely
 - and the negative pairs less likely,
 - over the entire training set



Word2vec

- Skip-gram models learns two sets of embeddings
 - Target embeddings matrix W
 - Context embedding matrix C
- It's common to just add them together, representing word i as the vector $w_i + c_i$

Summarize Word2vec

- Start with V random d -dimensional vectors as initial embeddings
- Train a classifier based on embedding similarity
 - Take a corpus and take pairs of words that co-occur as positive examples
 - Take pairs of words that don't co-occur as negative examples
 - Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
- Throw away the classifier code and keep the embeddings.

Word2vec – SGD

$$L_{CE} = - \left[\log \sigma(c_{pos} \cdot w) + \sum_{i=1}^k \log \sigma(-c_{neg_i} \cdot w) \right]$$

$$\frac{\partial L_{CE}}{\partial c_{pos}} = [\sigma(c_{pos} \cdot w) - 1]w$$

$$\frac{\partial L_{CE}}{\partial c_{neg}} = [\sigma(c_{neg} \cdot w)]w$$

$$\frac{\partial L_{CE}}{\partial w} = [\sigma(c_{pos} \cdot w) - 1]c_{pos} + \sum_{i=1}^k [\sigma(c_{neg_i} \cdot w)]c_{neg_i}$$

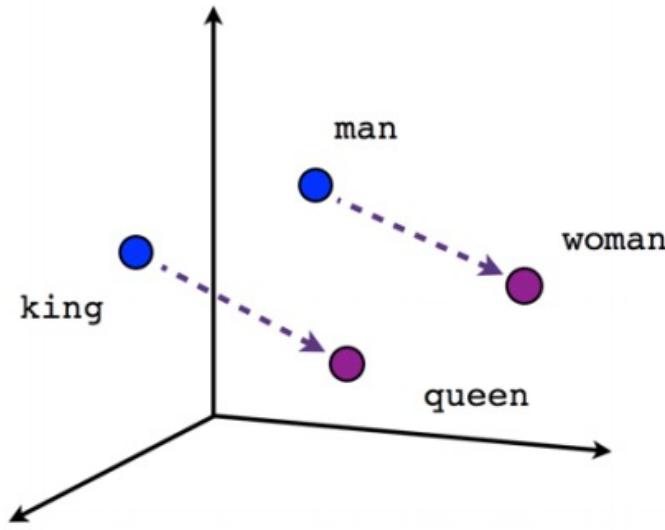
Update equation in SGD: Start with randomly initialised C and W matrices, then incrementally do updates

$$c_{pos}^{t+1} = c_{pos}^t - \eta [\sigma(c_{pos}^t \cdot w) - 1]w$$

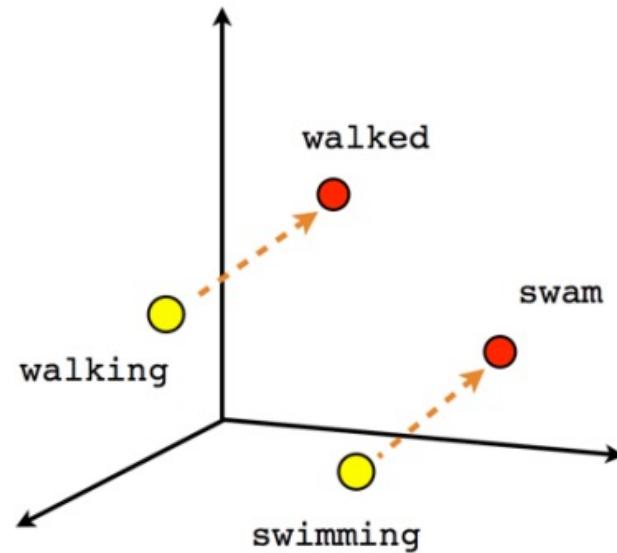
$$c_{neg}^{t+1} = c_{neg}^t - \eta [\sigma(c_{neg}^t \cdot w)]w$$

$$w^{t+1} = w^t - \eta [\sigma(c_{pos} \cdot w^t) - 1]c_{pos} + \sum_{i=1}^k [\sigma(c_{neg_i} \cdot w^t)]c_{neg_i}$$

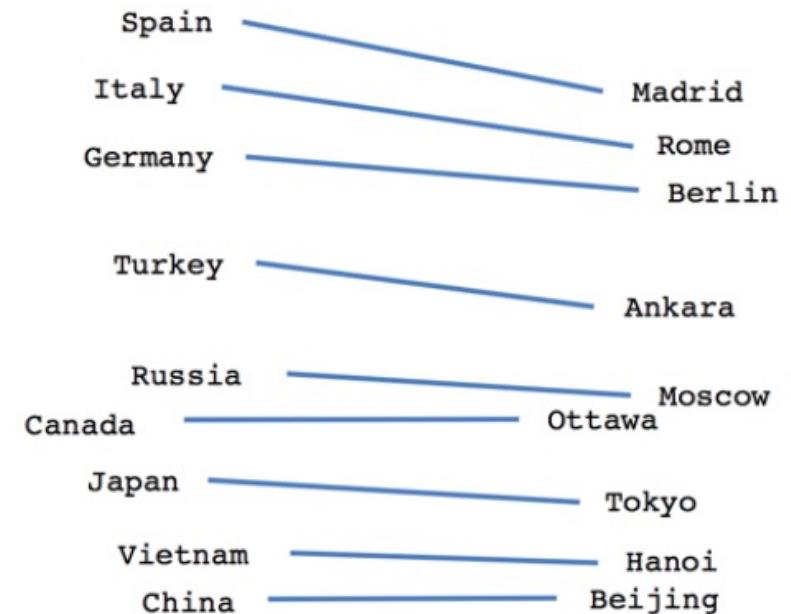
How does the vector space look like?



Male-Female



Verb tense



Country-Capital

Next lecture – do you know?

- Stochastic Gradient Descent ?
- Hierarchical Softmax ?

Literature

- Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, *Introduction to Information Retrieval*, Cambridge University Press. 2008.
- <https://stanford-cs324.github.io/winter2022/lectures/introduction/>
- https://stanford.edu/~jurafsky/slp3/slides/6_Vector_Jan13_2020.pdf
- Video Lectures of Andrew Ng.
- Papers:
 - [Efficient Estimation of Word Representations in Vector Space](#)
 - [Distributed Representations of Words and Phrases and their Compositionality](#)